

ctc-e-branchformer-asr

Problem statement

This project implements a CTC-only Automatic Speech Recognition (ASR) system using the E-Branchformer encoder, trained on the LibriSpeech 100-hour dataset. The model integrates a Language Model (LM) during inference to enhance decoding accuracy and reduce Word Error Rate (WER).

Repository structure

```
ctc-e-branchformer-asr/  
|-- requirements.txt  
|-- data/  
|   |-- (place LibriSpeech files or links here)  
|-- scripts/  
|   |-- build_lm.sh  
|   |-- prepare_manifest.py  
|-- src/  
|   |-- features.py  
|   |-- model.py  
|   |-- train.py  
|   |-- decode.py  
|   |-- evaluate.py  
|-- configs/  
|   |-- config.yaml  
|-- assets/  
    |-- (place training curves / sample outputs)
```

Requirements

```
torch>=1.12  
torchaudio  
numpy  
scipy  
librosa  
pyctcdecode  
kenlm  
jiwer  
tqdm  
PyYAML  
soundfile
```

Example configuration (configs/config.yaml)

```
dataset:  
  sample_rate: 16000  
  n_mels: 80  
  n_fft: 512  
  win_length: 400  
  hop_length: 160
```

```
train:
  batch_size: 16
  epochs: 60
  lr: 1e-4
  device: cuda

model:
  input_dim: 80
  d_model: 256
  num_layers: 12
  nhead: 4
  conv_expansion: 2
  dropout: 0.1

decode:
  beam_width: 100
  lm_weight: 2.0
  word_score: -1.0
```

Quick start (minimal instructions)

1. Install dependencies:
`pip install -r requirements.txt`
2. Prepare manifests (example):
`python scripts/prepare_manifest.py --librispeech_root /path/to/train-clean-100 --out data/train_manifest.json`
3. (Optional) Build LM:
`jq -r '.text' data/train_manifest.json > data/train_transcripts.txt`
`./scripts/build_lm.sh data/train_transcripts.txt models/kenlm.arpa`
4. Train:
`python src/train.py --manifest data/train_manifest.json --config configs/config.yaml --out checkpoints`
5. Decode / evaluate:
`python src/decode.py --wav sample.wav --ckpt checkpoints/checkpoint_epoch10.pt --lm models/kenlm.arpa`
`python src/evaluate.py --manifest data/test_manifest.json --ckpt checkpoints/checkpoint_epoch50.pt --lm models/kenlm.arpa`

Notes / tips

- Ensure LibriSpeech files are placed under data/ and manifests point to correct paths.
- If you get OOM errors, reduce batch_size or num_layers.
- Building KenLM (Implz) requires system build tools; use Docker or follow KenLM instructions if building fails.
- Keep feature extraction parameters consistent across train/val/test.

Deliverables

- Training logs and validation curves (place under assets/)
- Model explanation and configuration (configs/config.yaml, src/model.py)
- WER comparison (with and without LM) — include results file in assets/